

## Searching

Organizing and retrieving information is at the heart of most computer applications

- Searching is a very frequently performed task
- Search process:
  - ☐ Abstract view: Determine if an element with a particular value is a member of a particular set
  - ☐ Common view: Try to find the record with a record collection that has a particular key value.
- Some of the techniques presented here require material from chapter 8.
  - ☐ Assigned reading: Chapter 8, section 8.3 and all of Chapter 9

## Buffers and Buffer Pools (sec 8.3)

The general idea is to use a RAM buffer to hide latency.

- **Caching** or **buffering**: the act of storing in RAM a piece of data from a faster or slower device
  - ☐ allows the faster device to do something else while the slower device reads from or writes to the buffer
- Examples
  - ☐ CPU cache is a buffer for RAM
  - ☐ RAM is a buffer for disks of various types
  - ☐ Disk can buffer for tape
- Associated concepts:
  - ☐ **Buffer pool**: a set of multiple buffers
  - ☐ **Page**: a piece of memory large enough to fill a buffer

## Buffer Pools

---

An example is virtual memory.

- A hard disk is used to simulate a very large RAM memory
- System RAM is the buffer pool
- A **page** is a block of memory (usually some multiple of 512 bytes)
  - Address space of a process can be broken into multiple pages
  - At any given time, some pages may be on disk and some in RAM
    - Requesting a memory address currently on disk causes a **page fault**
  - Two options for page fault:
    - Find an “empty” page in RAM and transfer the page from disk
    - No empty pages in RAM: follow a page replacement strategy
  - Page replacement strategies:
    - FIFO
    - LFU
    - LRU

## Searching

---

- Formal definition:
  - Suppose  $k_1, k_2, \dots, k_n$  are distinct keys
  - Given a collection  $C$  of  $n$  records of the form
    - $(k_1, I_1), (k_2, I_2), \dots, (k_n, I_n)$
  - $I_j$  is information associated with key  $k_j$  for  $1 \leq j \leq n$ .
- **Search problem:** given key value  $K$ , locate the record  $(k_j, I_j)$  in  $C$  such that  $k_j = K$ 
  - **successful** search: record with  $k_j = K$  is found
  - **unsuccessful** search: no record with  $k_j = K$  is found
- Queries:
  - **Exact-match query:** search for a record whose key matches a specific key value
  - **Range query:** search for all records whose key values fall within a specified range

## Searching Categorization

---

- Three general approaches
  - ☐ Sequential and list methods
    - Works well for sequences (duplicate keys allowed)
    - Appropriate for data stored in RAM
  - ☐ Direct access by key value (hashing)
    - Doesn't work well for sequences
    - Works well for data on disk or in RAM
  - ☐ Tree indexing methods (chapter 10, not covered)

## Searching Sorted Arrays

---

- Sequential search
  - ☐  $\Theta(n)$ , average and worst case
  - ☐ Unacceptable for large data sets
- Binary search
  - ☐  $\Theta(\log n)$ , average and worst case
  - ☐ Works only for previously sorted data
- Dictionary search
  - ☐ a “computed” binary search
  - ☐ based on knowledge about key distribution
  - ☐ also called interpolation search

## Lists Ordered by Frequency

Instead of ordering by key value, a list may be ordered by frequency of access.

- Lists ordered by frequency: the *expected* frequency of occurrence determines ordering strategy

☐ A sequential search is performed

- Cost to access  $i^{\text{th}}$  record is  $i$
- Order in decreasing order of probability:  $p_i$  is the probability that record  $i$  will be accessed
- That is,

$$p_1 \geq p_2 \geq \dots p_n$$

(Note:  $\sum_{i=1}^n p_i = 1$  must be true)

- The cost to access each element is (position of element)  $\times$  (probability of element)
- Then the overall expected search cost is

$$\overline{C}_n = 1p_1 + 2p_2 + \dots + np_n$$

## Lists Ordered by Frequency (cont.)

- Example: all records have equal probability

☐  $p_i = 1/n$

☐ Then

$$\overline{C}_n = 1 \times 1/n + 2 \times 1/n + \dots + n \times 1/n$$

$$= \sum_{i=1}^n i/n = \frac{1}{n} \sum_{i=1}^n i$$

$$= \frac{1}{n} \times \frac{n(n+1)}{2} = \frac{n+1}{2}$$

- Example: exponential frequency

☐ Probabilities:

$$p_i = \begin{cases} 1/2^i & \text{if } 1 \leq i \leq n-1 \\ 1/2^{n-1} & \text{if } i = n \end{cases}$$

☐ Thus,

$$\overline{C}_n \approx \sum_{i=1}^n \frac{i}{2^i} \approx 2$$

## The 80/20 Rule

Many real access patterns follow this rule of thumb.

- The **80/20 rule**: 80
  - 80 and 20 are estimates (applications have their own values)
  - This behavior justifies caching techniques
  - When the rule applies, then reasonable search performance can be expected
- Example: **Zipf distribution**
  - A pattern followed by some naturally occurring distributions, including:
    - Distribution for frequency of word usage
    - Distribution for city populations
  - Related to the Harmonic series (chapter 2) as follows:
    - Zipf frequency for item  $i$  is  $1/i\mathcal{H}_n$
    - (here  $\mathcal{H}_n = \sum_{i=1}^n 1/i \approx \log_e n$ )
    - Then

$$\begin{aligned}\overline{C}_n &= \sum_{i=1}^n i/i\mathcal{H}_n \\ &= n/\mathcal{H}_n \\ &\approx n/\log_e n\end{aligned}$$

## Self-Organizing Lists

This is why we studied the section on buffer pools.

- A **self organizing list** is a list that starts out unordered, but the access policy includes procedures to impose an order based on actual pattern of record access
  - Use rules called **heuristics** to determine how to reorder the list
  - The heuristics are similar to the buffer pool management strategies (buffer pools are like a form of self-organizing list)
  - Heuristics:
    - **Count**: Count the frequency of access. When a record is found, increment its count and move it up if the count is greater than preceding record(s)
    - **Move-to-front**: when a record is found, move it to the front of the list
    - **Transpose**: when a record is found, swap it with the record ahead of it

## Self-Organizing Lists, Examples

---

- Initial list is A, B, C, D, E, F, G, H
- Access pattern is F D F G E G F A D F G E
  - ☐ Count heuristic:

☐ Move-to-front heuristic:

☐ Transpose heuristic:

## Self-Organizing Lists, Examples

---

- Application: text compression
  - ☐ Keep a table of words previously seen
  - ☐ Use the move-to-front heuristic
  - ☐ If a word is not yet seen, then send the word
  - ☐ If a word has been seen, then send its current table index

☐ Example: The car on the left hit the car I left

☐ becomes: The car on 3 left hit 3 5 I 5

☐ Similar in spirit to Ziv-Lempel coding

## Searching in Sets

---

Determining whether a value is a member of a set is a special case of searching for keys in a sequence of records.

- Any of the prior search methods can be used
- This problem allows us to speed up the process:
  - **Bit vector** or **bitmap** representation: use an array of  $n$  bits corresponding to  $n$  potential set members
    - $i = 1$  means that member  $i$  is present
    - $i = 0$  means that member  $i$  is not present
  - Application: document retrieval: find all documents in a set containing certain keywords
    - For each keyword, the system stores a bit vector (one bit for each document)
    - A '1' means that the document contains the keyword
    - Searching for three words is a logical AND of 3 bit vectors.

## Hashing

---

A completely different approach in which search is by direct access based on the key value.

- **Hashing** is the process of accessing a record by mapping a key value to a position in a table.
- The mapping process requires a (normally  $\Theta(1)$ ) mathematical function called the **hash function**, denoted by **h**
- The **Hash table** is an array that stores all of the records, denoted **HT**
- A record's position in the hash table is its **slot**
- The number of slots is denoted by  $M$ , numbering is from 0 to  $M - 1$
- The mapping function **h** must work as follows:
  - For any value  $K$  in the key range,  
 $h(K) = i, 0 \leq i < M$   
such that  $\text{key}(\mathbf{HT}(i)) = K$

## Hashing (cont.)

---

Hashing answers the specific question “what record, if any, has key value  $K$ ?”

- Works well for sets (no duplicates)
- Not suitable for range queries
- Works well for in-memory and disk-based applications
- Example:
  - ☐ Store the  $n$  records with key values in the range 0 to  $n - 1$
  - ☐ Hash function  $h(K) = K$
  - ☐ This is not a practical example (Why?)
- Example:
  - ☐ Store about 1000 records having keys in the range 0 to 16,383
  - ☐ Impractical to keep a hash table with 16,383 slots
  - ☐ We need a hash function that maps the key range to a smaller table

## Collisions

---

- Given a hash function  $h(k)$  and keys  $k_1$  and  $k_2$ :
  - ☐ If  $h(k_1) = h(k_2) = \beta$ , then  $k_1$  and  $k_2$  have a **collision** at  $\beta$  under  $h$ .
- Collisions are inevitable in most applications
  - ☐ Example: birthday sharing
- Minimizing collisions requires good hash functions
- Finding a record (or a place in which to insert) requires a two-step procedure:
  1. Compute table location  $h(k)$
  2. Starting with slot  $h(k)$ , search for the record containing key  $k$  (or an empty location where it may be inserted)
- The search procedure is the collision resolution technique. There are two major classes:
  - ☐ **Open hashing**, also called **Separate chaining**
  - ☐ **Closed hashing**, also called **Open addressing**



## Hash Functions

- Requirement:
  - ☐ A hash function must compute a slot index within the hash table's range; thus, it computes  $(\text{some value}) \bmod M$
- Goals:
  - ☐ A practical hash function evenly distributes the records stored among the hash table slots
  - ☐ Ideally, the even distribution is to all slots with equal probability
    - Success at this depends on the data's distribution
  - ☐ It should also be fast (probably the easiest goal to accomplish)
- Two situations normally faced:
  - ☐ We know nothing about the incoming key distribution: attempt to evenly distribute the key range over the hash table, trying to avoid clustering
  - ☐ We know something about the incoming key distribution: use a distribution-dependent hash function.

## Examples

- A Simple hash function:

```
int h (int x) {
    return (x % 16);
}
```

  - ☐ The  $\bmod 16$  operation makes the function dependent on the lower 4 bits of the key
- Mid-square method: square the key value, taking the middle  $r$  bits from the result for a hash table having  $2^r$  slots
- Folding method: sum the ASCII values of all letters, taking the result  $\bmod M$ :

```
int h(char *x) {
    int i = 0; int sum = 0;
    while (x[i] != NULL) {
        sum += (int) x[i];
        i++;
    }
    return (sum % M);
}
```

## Examples

- Executable and Linking Format (ELF) hash, Unix Sys/V Release 4:

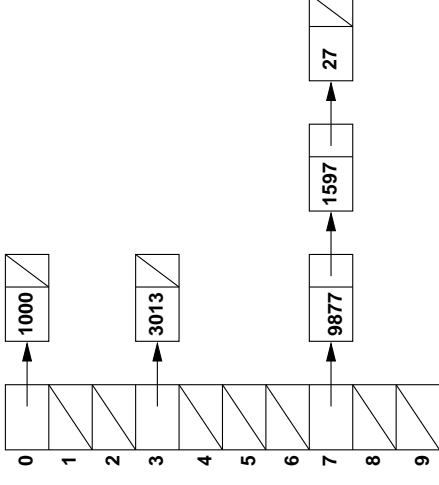
```
int ELFHash(char *key) {
    unsigned long h = 0;
    while (*key) {
        h = (h << 4) + *key++;
        unsigned long g = h & 0xF0000000L;
        if (g) h ^= g >> 24;
        h &= ~g;
    }
    return h % M;
}
```

- ☐ Works well with short and long strings
- ☐ Every letter of the string has equal effect
- ☐ Even distribution into hash table slots is very likely

## Open Hashing

This is also called **separate chaining**.

- A collision resolution technique, in which:
  - ☐ The hash table is not an array of records; rather, it is an array of pointers
  - ☐ Each slot is treated as a bin so that collisions do not really occur
  - ☐ For a given record with key  $k$  and  $h(k) = \beta$ :
    - hash table slot  $\beta$  is the head of a linked list
    - Insert into slot  $\beta$  becomes a linked list insert



## Closed Hashing

This is also called **open addressing**

- All records are stored directly in the hash table
  - Each record  $i$  has a **home position** defined by  $h(k_i)$ 
    - If record  $i$  is inserted and another record already occupies  $i$ 's home position, then another slot must be found to store  $i$ .
    - The search procedure to find a new slot is the **collision resolution policy**

## Bucket Hashing

One implementation of closed hashing in which the extra list space is stored in the table.

- Divide the hash table into buckets
  - $M$  slots are divided into  $B$  buckets, with  $M > B$
  - Include overflow bucket with large capacity at end
  - Records hash to the first slot of the bucket, then fill it sequentially
  - Overflow is used if a given bucket is full
  - Search: check bucket then check overflow (using linear search in both)

## Collision Resolution Policies

- Goal is to find a free slot in the table
- Search proceeds by following a **probe sequence**: the series of slots visited during insert/search after a collision occurs
  - ☐ Whether inserting or searching, the probe sequence must be the same every time
  - ☐ Basic idea: follow probe sequence until one of the following is true:
    - record with key =  $k$  is found
    - an empty slot is found (no record with key  $k$  exists in the hash table)

- ☐ Insert with Probing:

```
void insert(item R) {
    int home, pos, i;
    home = h(key(R));
    if (Table[home] == EMPTY)
        Table[home] = R;
    else {
        for (i = 1; Table[pos] != EMPTY; i++) {
            pos = (home + probe(key(R), i)) % M;
            if (key(T[pos]) == key(R)) ERROR;
        }
        Table[pos] = R;
    }
}
```

## Linear Probing

From a given position, linear probing searches the next available slot in the table.

- Probe function:

```
int probe(int Key, int i) { return i; }
```

- ☐ If the end of the table is reached, it wraps around to the top (see code on previous page)
- ☐ At least one slot must always be empty in the table. Why?

- Linear probing suffers from **primary clustering**:

- ☐ “Clusters” of occupied cells form
- ☐ Any key hashing into a cluster requires several attempts to resolve the collision and then will add to the cluster

## Primary Clustering

- Probabilities for which slot to use next are not the same
  - ☐  $h(k) = k \bmod 11$
  - ☐  $1003 \bmod 11 = 2, 1924 \bmod 11 = 10, 3071 \bmod 11 = 2, 2071 \bmod 11 = 3, 4752 \bmod 11 = 0$

Insert in the following order: 1003, 1924, 3071, 2071, 4752

0	4752
1	
2	1003
3	3071
4	2071
5	
6	
7	
8	
9	
10	1924

## Better Linear Probing

- Use a constant  $c$  to skip by, instead of going to the next slot on every probe
  - ☐  $\text{probe}(h(k), i) = h(k) + c \times i$
  - ☐  $M$  and  $c$  should be relatively prime (Why?)
- Clustering can still exist
  - ☐ Example:  $c = 3, h(k_1) = 3, h(k_2) = 9$
  - ☐ Probe sequences for  $k_1$  and  $k_2$  are linked together

## Pseudo Random Probing

---

An ideal probe function selects the next slot in the probe sequence at random

- Why can a real probe function not act randomly?
- Pseudo random probing:
  - ☐ Select a random permutation of the numbers from 1 to  $M - 1$ :  $r_1, r_2, \dots, r_{M-1}$
  - ☐ All searches and insertions use the same permutation:
  - ☐  $p(K, i) = \text{Perm}[i - 1]$
  - ☐ that is, the  $i^{\text{th}}$  value in the probe sequence is  $(h(k) + r_i) \bmod M$
- Example:
  - ☐  $M = 101$
  - ☐  $r_1 = 2, r_2 = 5, r_3 = 32$
  - ☐  $h(k_1) = 30, h(k_2) = 28$
  - ☐ Probe sequence for  $k_1$ :
  - ☐ Probe sequence for  $k_2$ :

## Quadratic Probing

---

- The  $i^{\text{th}}$  probe sequence function is  $i^2$
- That is, the  $i^{\text{th}}$  value in the probe sequence is  $(h(k) + i^2) \bmod M$
- Example:
  - ☐  $M = 101$
  - ☐  $h(k_1) = 23, h(k_2) = 24$
  - ☐ Probe sequence for  $k_1$ :
  - ☐ Probe sequence for  $k_2$ :


## Double Hashing

Prior probing methods can reduce or eliminate primary clustering.

- **Secondary clustering** occurs when two keys hash to the same slot, thus following the exact same probe sequence
- Desirable: the probe sequence is a function of both the key and the home position
- Double hashing adds a second hash function to the probe sequence:
  - ☐  $p(k, i) = i \times h_2(k)$  for  $0 \leq i \leq M - 1$
  - ☐ Poor choice of  $h_2(k)$  results in poor ("disastrous") performance
  - ☐ Make sure all cells can be probed by ensuring that all probe sequence constants are relatively prime to  $M$ 
    - One method: make  $M$  prime
    - Another method: set  $M = 2^m$  and make  $h_2$  return an odd value between 1 and  $2^{m-1}$

## Analysis of Closed Hashing

- Visualizing the expected performance of hashing based on load factor
- Load factor  $\alpha = N/M$  where  $N$  is the number of records stored



- Note: "random probe" is only a theoretical measure

## Rehashing

---

- Consequences of a hash table that is too full:
  - ☐ Running time for operations start to take too long
  - ☐ Insertions might fail for certain collision resolution strategies
- Solution: build a bigger table
  - ☐ Find a prime number at least twice as large as current value of  $M$
  - ☐ Allocate a new hash table (array)
  - ☐ Scan through the old hash table, inserting all elements into the new hash table
  - ☐ Delete the old hash table
- Operation is expensive but occurs relatively infrequently
- Strategies:
  - ☐ Rehash when the table is half full
  - ☐ Rehash when an insertion fails
  - ☐ Rehash when the table reaches a certain **load factor**

## Deletion

---

Deletion is tricky with hashing for the following reasons:

- Deleting a record must not hinder later searches (an empty slot means “stop the search”)
- Positions should also not be made unusable due to deletions (avoid a “zombie slot?”)
- Solution:
  - ☐ Add a special mark in place of the deleted record.
  - ☐ Mark is called the **tombstone**
  - ☐ Tombstones do not stop search but do add to average search time
  - ☐ Solutions to that added time:
    - Local reorganizations to try to shorten it
    - Periodically rehash the table